

Probabilistic Model Checking

A few hints on probabilities in verification

Absolute guarantee of correctness?


- Often system are subject to stochastic phenomena
 - Message loss, hardware failure, ...
- They can initially be tackled with nondeterminism, but we'd like to have some bounds on probability
 - "If Messages can be lost 1% of the times then correct answers are delivered 90% of the times"
- For this, we need to extend our models to allow probabilities
 - Markov Chains, Markov Decision Processes
- And to verify their properties
 - Qualitative: "good eventually happens with probability 1", "bad happens with probability 0"
 - Quantitative: "Good eventually happens with at least 95% probability", "Bad happens with less than 5% probability"

Discrete-time Markov Chains

- Markov chains are the most popular operational model for the evaluation of *performance* and *dependability* of information-processing systems.
- Roughly speaking, Markov chains are transition systems with probability distributions for the successors of each state.
 - That is, instead of a nondeterministic choice, the next state is chosen probabilistically.
- They lack nondeterminism, hence they cannot model interleaving behavior of concurrent processes

An example: lost messages

In state *try*, a message is delivered, but it is lost with probability $1/10$, in which case the message will be sent off again, until it is eventually delivered.



Formal Definition

A *(discrete-time) Markov chain* is a tuple $\mathcal{M} = (S, \mathbf{P}, \iota_{\text{init}}, AP, L)$ where


- S is a countable, nonempty set of states,
- $\mathbf{P} : S \times S \rightarrow [0, 1]$ is the *transition probability function* such that for all states s :

$$\sum_{s' \in S} \mathbf{P}(s, s') = 1,$$

- $\iota_{\text{init}} : S \rightarrow [0, 1]$ is the *initial distribution*, such that $\sum_{s \in S} \iota_{\text{init}}(s) = 1$, and
- AP is a set of atomic propositions and $L : S \rightarrow 2^{AP}$ a labeling function.

Example: transition probability matrix

Using the enumeration start, try, lost, delivered for the states, the transition probability function P can be viewed as a 4×4 matrix and the initial distribution as a column vector



$$\mathbf{P} = \begin{pmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & \frac{1}{10} & \frac{9}{10} \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{pmatrix} \quad l_{\text{init}} = \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

An example of a path is $\pi = (\textit{start try lost try lost try delivered})^\omega$

Along this path each message has to be retransmitted two times before delivery. It follows that $\inf(\pi) = S$. For $T = \{\textit{lost}, \textit{delivered}\}$, we have $\mathbf{P}(\textit{try}, T) = 1$. ■

LTL? CTL?

- By abstracting away probabilities, a DTMC is a transition system TS, obtained by considering non-zero probability initial states and transitions. Can we use LTL/CTL model checking?
- Yes, but in the example the following CTL formula holds:

$$start \models \exists \square \neg delivered$$

- Because we can build an infinite path where a message is never delivered! But it is clear that this path has “0 probability”
- So we need a different logic!

LTL-style formulae and reachability probabilities

- We can use LTL to talk about events, e.g. $GF\ B$ describes the event that the subset B of states is visited infinitely often.
- Given a LTL formula φ the probability that φ holds in state s is:

$$Pr(s \models \varphi) = Pr_s\{ \pi \in Paths(s) \mid \pi \models \varphi \}.$$


- Pr_s is the total probability of the sets of paths with initial state s , for which the path formula φ holds.

Example: the lossy channel M

- Can we reach the state «delivered»?
- We are interested in path fragments $s_0...s_n$ with $s_n = delivered$ and $s_i \neq delivered$ for $0 \leq i < n$

$$\widehat{\pi}_n = start\ try\ (lost\ try)^n\ delivered$$

- Probability of $\widehat{\pi}_n$ is $\left(\frac{1}{10}\right)^n \cdot \frac{9}{10}$



$$Pr^{\mathcal{M}}(\Diamond delivered) = \sum_{n=0}^{\infty} \left(\frac{1}{10}\right)^n \cdot \frac{9}{10} = \frac{\frac{9}{10}}{1 - \frac{1}{10}} = \frac{\frac{9}{10}}{\frac{9}{10}} = 1.$$

PCTL: Probabilistic Computation Tree Logic

- A branching time temporal logic based on CTL

- Syntax

$$\Phi ::= \text{true} \mid a \mid \Phi_1 \wedge \Phi_2 \mid \neg \Phi \mid \mathbb{P}_J(\varphi)$$

where $a \in AP$, φ is a path formula and $J \subseteq [0, 1]$ is an interval with rational bounds. PCTL *path formulae* are formed according to the following grammar:

$$\varphi ::= \bigcirc \Phi \mid \Phi_1 \cup \Phi_2 \mid \Phi_1 \cup^{\leq n} \Phi_2$$

where Φ , Φ_1 , and Φ_2 are state formulae and $n \in \mathbb{N}$. ■

abbreviations are used; e.g., $\mathbb{P}_{\leq 0.5}(\varphi)$ denotes $\mathbb{P}_{[0,0.5]}(\varphi)$, $\mathbb{P}_{=1}(\varphi)$ stands for $\mathbb{P}_{[1,1]}(\varphi)$, and $\mathbb{P}_{>0}(\varphi)$ denotes $\mathbb{P}_{]0,1]}(\varphi)$.

Some intuition on non-probabilistic operators

The propositional logic fragment of PCTL, as well as the path formulae $\bigcirc \Phi$ and $\Phi_1 \cup \Phi_2$ has the same meaning as in CTL. Path formula $\Phi_1 \cup^{\leq n} \Phi_2$ is the *step-bounded* variant of $\Phi_1 \cup \Phi_2$. It asserts that the event specified by Φ_2 will hold within at most n steps, while Φ_1 holds in all states that are visited before a Φ_2 -state has been reached. Other Boolean connectives are derived in the usual way, e.g., $\Phi_1 \vee \Phi_2$ is obtained by $\neg(\neg\Phi_1 \wedge \neg\Phi_2)$. The eventually operator (\Diamond) can be derived as usual: $\Diamond\Phi = \text{true} \cup \Phi$. Similarly, for step-bounded eventually we have:

$$\Diamond^{\leq n} \Phi = \text{true} \cup^{\leq n} \Phi.$$

A path satisfies $\Diamond^{\leq n} \Phi$ if it reaches a Φ -state within n steps.

The always operator can be derived using the duality of eventually and always and the duality of lower and upper bounds.

$$\mathbb{P}_{\leq p}(\Box\Phi) = \mathbb{P}_{\geq 1-p}(\Diamond\neg\Phi) \quad \text{and} \quad \mathbb{P}_{]p,q]}(\Box^{\leq n}\Phi) = \mathbb{P}_{[1-q,1-p[}(\Diamond^{\leq n}\neg\Phi).$$

Usage of probabilistic operator

- Example: a 6-face die: 1 ,2,3,4 5 and 6 are the possible outcome

- If the die is fair, then:
$$\bigwedge_{1 \leq i \leq 6} \mathbb{P}_{=\frac{1}{6}} (\Diamond i)$$

- A communication protocol with a lossy channel

$$\mathbb{P}_{=1}(\Diamond delivered) \wedge \mathbb{P}_{=1} \left(\Box (try_to_send \rightarrow \mathbb{P}_{\geq 0.99}(\Diamond^{\leq 3} delivered)) \right)$$

Almost surely some message will be delivered (first conjunct) and that almost surely for any attempt to send a message, with probability at least 0.99, the message will be delivered within three steps.

PCTL semantics

- Semantics for PCTL over a Markov chain has same definition as CTL over a TS for the “non-probabilistic fragment”
- The semantics of the probability operator for a state s is:

$$s \models \mathbb{P}_J(\varphi) \quad \text{iff} \quad Pr(s \models \varphi) \in J.$$

Here, $Pr(s \models \varphi) = Pr_s\{ \pi \in Paths(s) \mid \pi \models \varphi \}$.

Pr_s is the total probability of the sets of paths with initial state s , for which the path formula φ holds.

PCTL semantics of the other operators

$$s \models a \quad \text{iff} \quad a \in L(s),$$

$$s \models \neg\Phi \quad \text{iff} \quad s \not\models \Phi,$$

$$s \models \Phi \wedge \Psi \quad \text{iff} \quad s \models \Phi \text{ and } s \models \Psi,$$

$$\pi \models \bigcirc \Phi \quad \text{iff} \quad \pi[1] \models \Phi,$$

$$\pi \models \Phi \cup \Psi \quad \text{iff} \quad \exists j \geq 0. (\pi[j] \models \Psi \wedge (\forall 0 \leq k < j. \pi[k] \models \Phi)),$$

$$\pi \models \Phi \cup^{\leq n} \Psi \quad \text{iff} \quad \exists 0 \leq j \leq n. (\pi[j] \models \Psi \wedge (\forall 0 \leq k < j. \pi[k] \models \Phi))$$

where for path $\pi = s_0 s_1 s_2 \dots$ and integer $i \geq 0$, $\pi[i]$ denotes the $(i+1)$ -st state of π , i.e., $\pi[i] = s_i$. ■

PCTL Model Checking

The PCTL model-checking problem is the following decision problem. Given a finite Markov chain \mathcal{M} , state s in \mathcal{M} , and PCTL state formula Φ , determine whether $s \models \Phi$. As for CTL model checking, the basic procedure is to compute the satisfaction set $Sat(\Phi)$. This is done recursively using a bottom-up traversal of the parse tree of Φ ; see Algorithm

Theorem 10.40. Time Complexity of PCTL Model Checking for MCs

For finite Markov chain \mathcal{M} and PCTL formula Φ , the PCTL model-checking problem $\mathcal{M} \models \Phi$ can be solved in time

$$\mathcal{O}(\text{poly}(\text{size}(\mathcal{M})) \cdot n_{\max} \cdot |\Phi|)$$

where n_{\max} is the maximal step bound that appears in a subpath formula $\Psi_1 \mathbf{U}^{\leq n} \Psi_2$ of Φ (and $n_{\max} = 1$ if Φ does not contain a step-bounded until operator).

Qualitative fragment of PCTL

- Allowing only “>0” “=1”, “=0”, “<1” as probability bound $\mathbb{P}_{=1}(a \cup b)$
 - Useful when we do not need to have explicit probability
 - Model Checking becomes more efficient
- This fragment is NOT equivalent to CTL, since, e.g, “with probability 0” can be different from “never”

There is no CTL formula that is equivalent to $\mathbb{P}_{=1}(\Diamond a)$.


There is no CTL formula that is equivalent to $\mathbb{P}_{>0}(\Box a)$.

There is no qualitative PCTL formula that is equivalent to $\forall \Diamond a$.

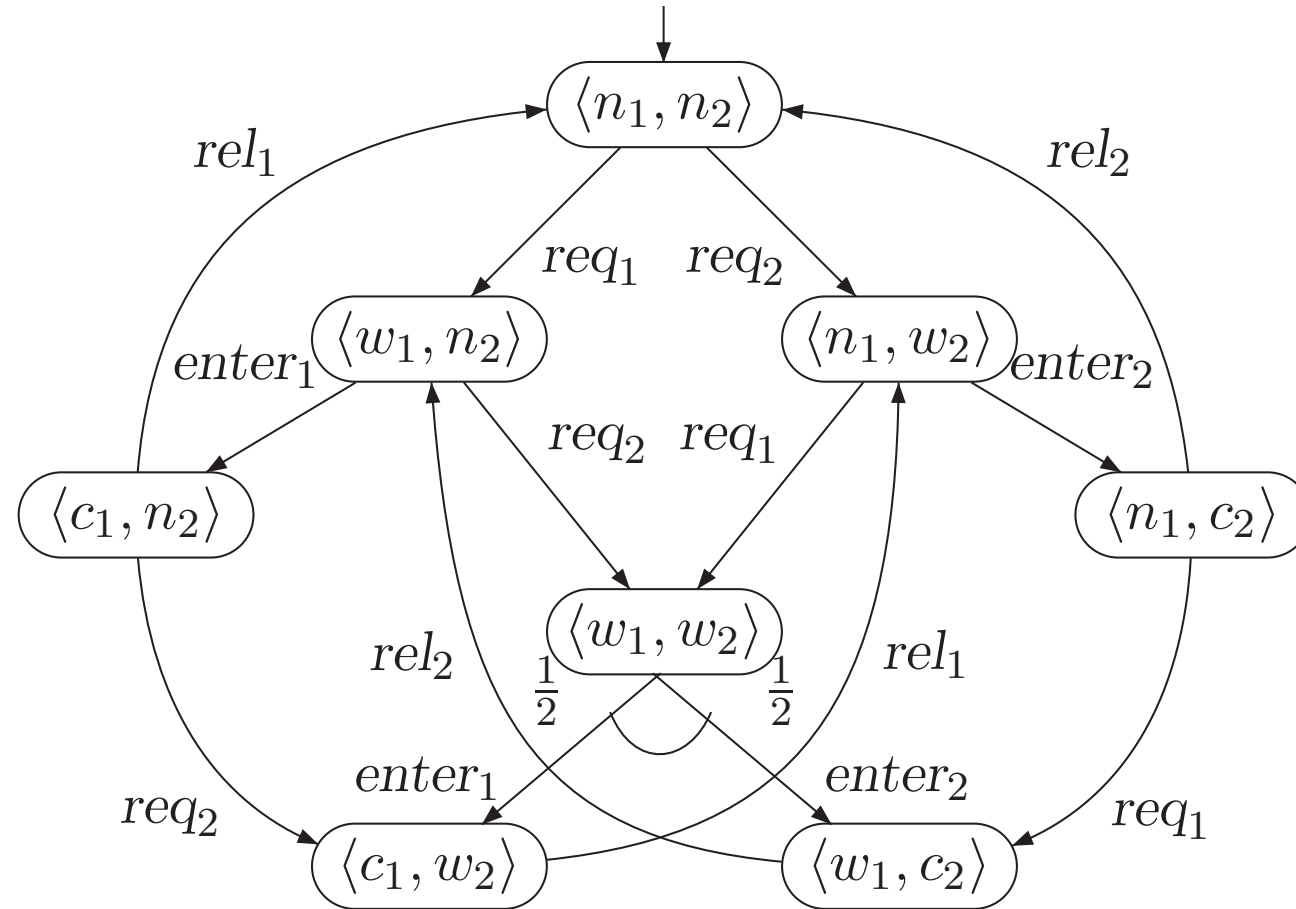
There is no qualitative PCTL formula that is equivalent to $\exists \Box a$.

Markov Decision Processes

- A variant of MC allowing nondeterminism
- Useful for modeling concurrent systems



A randomized mutual exclusion protocol



PCTL model checking over MDP

- Similar to the case for MC

Theorem 10.115. Time Complexity of PCTL Model Checking for MDPs

For finite MDP \mathcal{M} and PCTL formula Φ , the PCTL model-checking problem $\mathcal{M} \models \Phi$ can be determined in time

$$\mathcal{O}(\text{poly}(\text{size}(\mathcal{M})) \cdot n_{\max} \cdot |\Phi|)$$

where n_{\max} is the maximal step bound that appears in a sub-path formula $\Psi_1 \cup^{\leq n} \Psi_2$ of Φ (and $n_{\max} = 1$ if Φ does not contain a step-bounded until operator).